




# Effective Application Development with GemFire using Spring Data GemFire

John Blum – Luke Shannon

# Who Are We?




**John Blum**

*Spring Data GemFire Lead*

*GemFire Engineer/Tech Lead*

**Pivotal Software, Inc.**

**Twitter:** @john\_blum



**Luke Shannon**

*Community Engineer*

**Pivotal Software, Inc.**

# Agenda

- GemFire Overview
- Spring Data GemFire
  - Overview
  - Feature Enhancements
  - Productivity Features
  - New in 1.5.0.RELEASE
  - EA Concerns
  - Customer Case Studies
  - Looking Ahead
- Resources
- QA

# GemFire Overview

# What is GemFire?

GemFire is an *in-memory data grid* (IMDG)

GemFire can function as a OLTP/OLAP data management solution offering *high-throughput, low-latency* data access with *horizontal scale-out*

In the simplest case though, GemFire is just a *cache*...

- Key/Value Store (Objects/JSON)

# In-Memory

GemFire stores application data (state) in the **JVM Heap**

GemFire and application **performance** is impacted by GC...

- Object graphs can be stored in serialized form to reduce strain on the *Garbage Collector*.

No **off-heap** storage...

- GemFireXD, Pivotal's in-memory SQL-based store, supports off-heap

# In-Memory Continued...

Support for managing resources used by the local *cache* is handled by GemFire's *ResourceManager*

Important **configuration settings**:

- `critical-heap-percentage` - % of heap at which the *cache* might become inoperable due to GC pauses or `OutOfMemoryErrors`
- `eviction-heap-percentage` - % of heap when eviction begins on a Region configured with Heap LRU eviction

# Data Grid

GemFire pools system resources across multiple nodes joined in a cluster to manage the entire application's state and behavior...

- CPU, memory, network and (optionally) local disks

Members join the cluster as peers using either *UDP multicasting* or GemFire's *Locator* process...



# Shared-Nothing Architecture

Each member of the cluster is independent and self-sufficient...

- Members do not share memory or disk.

GemFire's "*Shared-Nothing Disk Persistence Architecture*" manages data in local disk files on each member independently from other members that host the same *cache (Region)*...

# High Scalability

Additional GemFire Server data nodes (peer members) can be added to the cluster, and data **rebalanced**, to uniformly distribute the load

- GemFire is “*elastic*”; easily scales up or down

GemFire really scales when data is **partitioned** across many members of the cluster uniformly...

- Application behavior can be routed, distributed and run in parallel with close proximity to the data required by application logic to improve **throughput**

# High-Throughput / Low-Latency

GemFire uses concurrent, in-memory data structures and a highly optimized distribution infrastructure to minimize context switching and contention

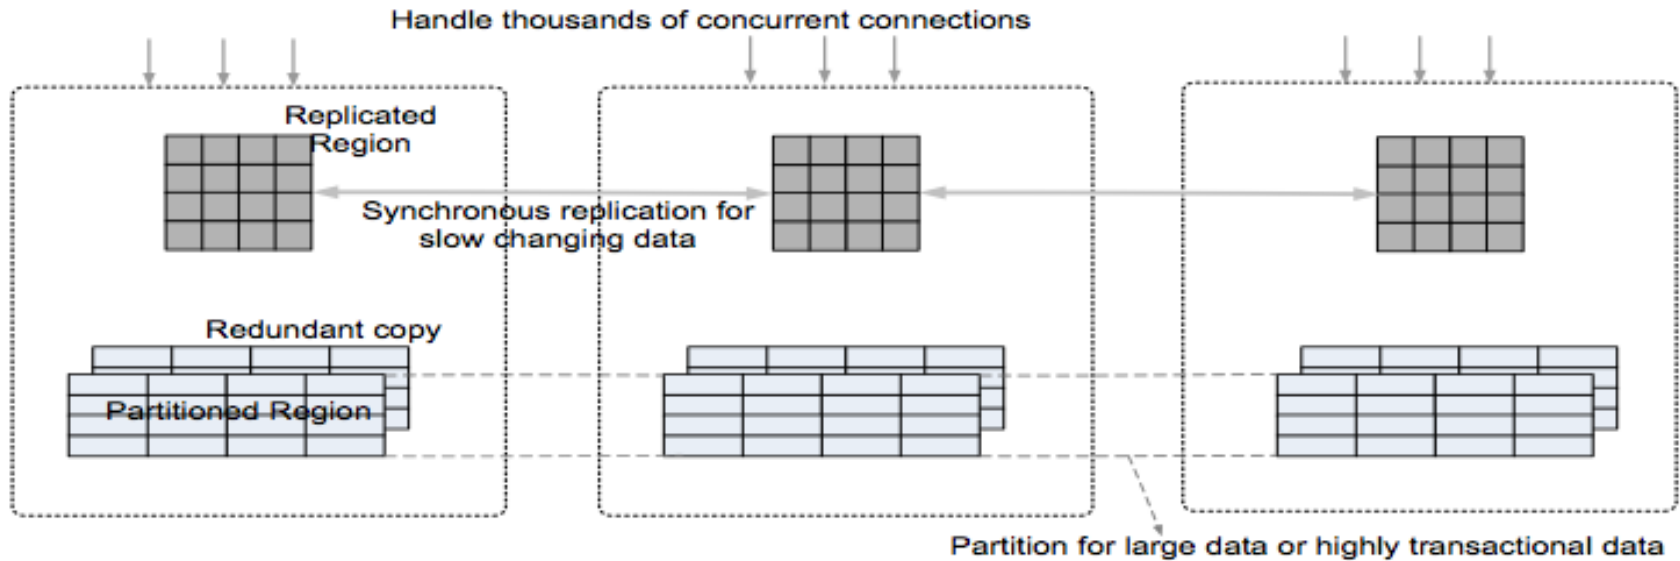
- **high read throughput** is achieved through *synchronous* or *asynchronous replication* of data, and...
- **high read and write throughput** is achieved when data is *partitioned* across many nodes in the cluster

*“Linear increase in throughput is limited only by the backbone network capacity.”*

# GemFire Data Storage Models

# REPLICATE vs. PARTITION

In GemFire, a *Region* is the *cache*...



# Region Configuration

*Cache-like* configuration settings...

- **Initial Capacity** and **Load Factor**
- **Key/Value-Constraints** (Class type)
- *Entry* **Eviction**
- *Region* and *Entry* **Expiration**

# Additional Region Configuration

Type – *DataPolicy, Shortcuts*

Persistence – *DataPolicy*

Distribution – *DataPolicy, Scope, Subscription*

Disk Store – *Overflow & Persistence*

Statistics (used for *Expiration*)

Async and Subscription Message Conflation

Async Event Queue (e.g. “*Write-Behind*”)

Gateway Sender (Site-to-Site WAN Replication)

# Partition Region Specific Settings


Local/Total Max Memory Size (MB)

Total Number of Buckets

Partition Resolver

Redundancy (# of Data Copies)

Collocation





# Other Region Types

## Distributed, Non-Replicated Regions

- Data is distributed; Each member only holds the data it has expressed “interests” in using *Subscription* or by defining cache entries (keys)

## Local Region

- Data is not distributed; local-only (`Scope.LOCAL`) only to the member defining the *Region*

## Empty Region

- Data accessor *Region* (PROXY); stores no data; receives events

## Client Region

- PROXY or CACHING\_PROXY

# Consistency

## Partition Regions

- *consistency* is maintained by routing all updates on a given key to the GemFire member holding the **primary** copy
- holds **lock** on key while distributing updates to other members having key

## Replicated Regions

- concurrent updates → possibility of **conflicts** and **out-of-order updates**
- *conflict checking* ensures all members eventually apply the same value (*eventual consistency*)
- *conflict checking* ensures out-of-order updates are discarded.

*Conflict checking* is enabled with... `--concurrency-checks-enabled`

# Cache Region Callbacks

## Cache Listener

- `afterCreate`, `afterUpdate`, `afterDestroy`, `afterInvalidate`, ...

## CacheLoader

- Synchronous “Read-Through” on *cache* misses from external data source

## CacheWriter

- Synchronous “Write-Through” to external data store

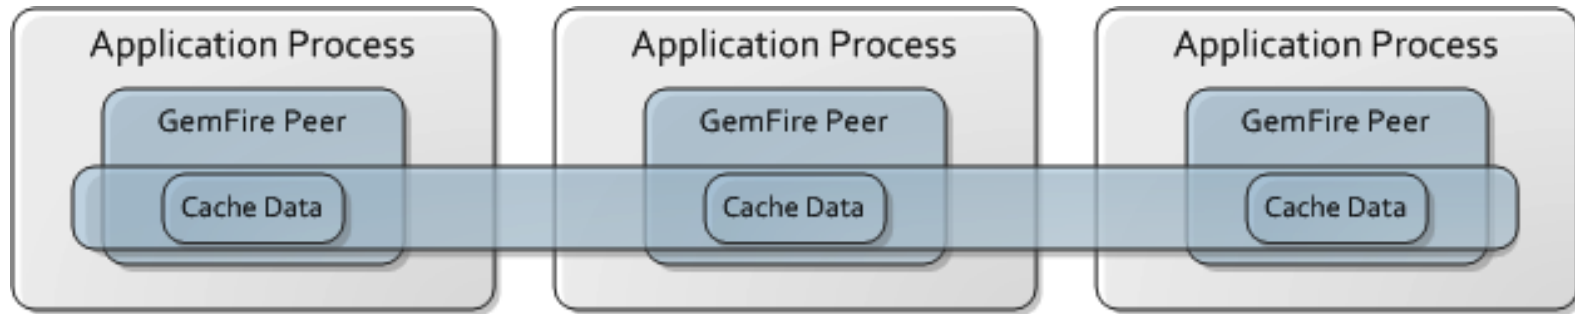
# GemFire Topologies

# Network Topologies

## Peer-to-Peer (P2P)

*Cache* configuration designed for applications embedding the *cache* within the same process space and participate in the cluster...

- close proximity of application logic and data



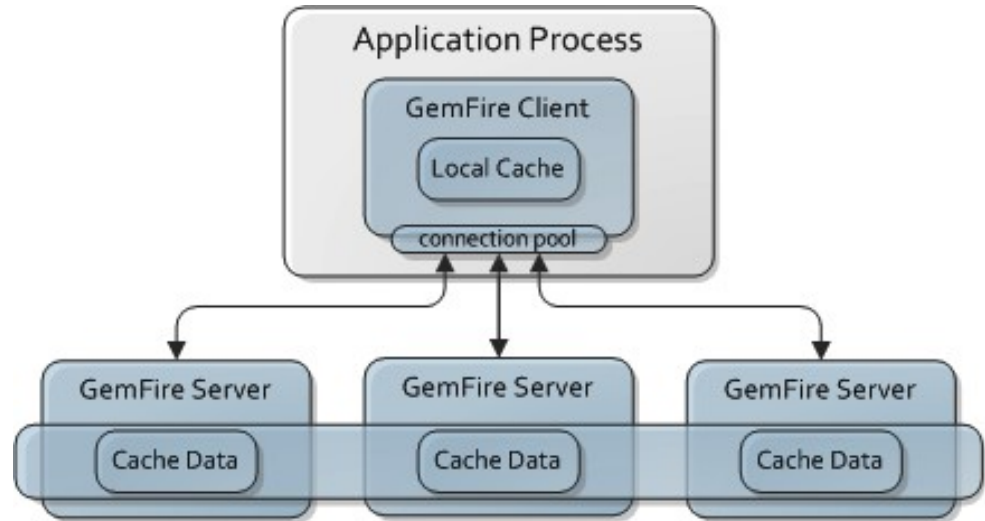
# Network Topologies

## Client / Server

Clients use single-hop data access...

Continuously load balanced between servers based on load to get more predictable response-times...

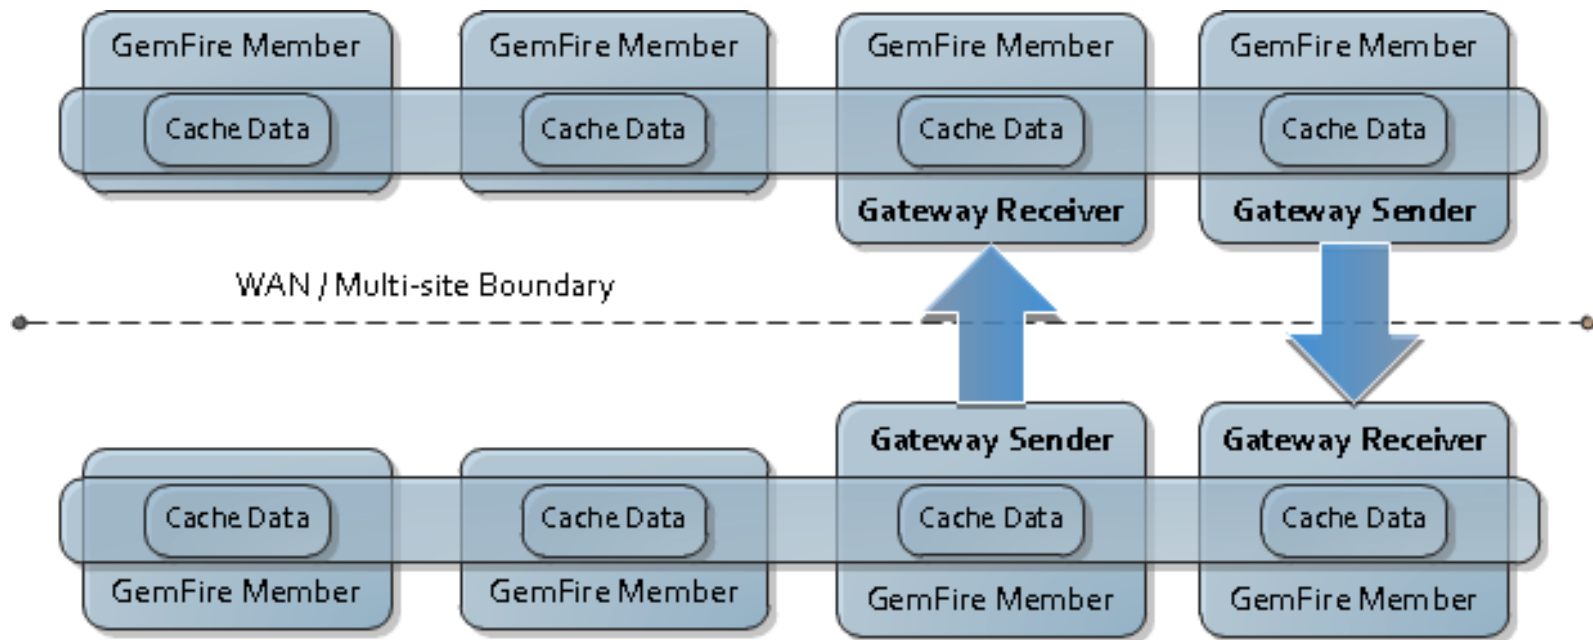
Automatic fail-over when Servers fail or become unresponsive



# Network Topologies

## WAN

### Site-to-Site Replication for DR



# GemFire Features++

Still just scratching the surface...

- Cache and JTA-compliant Transaction Management
- Functions
- OQL
- Client Register Interests or CQ
- Serialization (PDX, Java, DataSerialization)
- Security
- Native Support for Java and C++/C# Clients
- Memcache Integration
- REST Interface + JSON document storage
- Delta GII – when member **re-joins** the cluster
- Delta Propagation – distribute object changes only
- Rolling Upgrade - update GemFire version without having to bring the member/cluster down
- Split Brain Detection



# GemFire Tips

1. Do not directly modify cached values...
  - Bypasses GemFire distribution as well as cache listeners/writers, expiration activities, transaction management leading to undesired results
  - Instead, make a copy of the data by setting *–copy-on-read*, or use *CopyHelper*
2. Key class types must implement `Object equals(:Object)` and `hashCode()`;
3. GemFire serializes data entry keys and values during distribution, between client and server, and for persistence/overflow to Disk Stores, therefore keys/values must be serializable.
4. Application classes must be added to each peer member's CLASSPATH using the data in non-serialized form
5. Close the Region to release the resources held
6. Distributed System and Region configuration settings must match to be compatible systems

# GemFire Cluster Demo

# Spring Data GemFire

# Agenda

- GemFire Overview
- **Spring Data GemFire**
  - Overview
  - Feature Enhancements
  - Productivity Features
  - New in 1.5.0.RELEASE
  - EA Concerns
  - Customer Case Studies
  - Looking Ahead
- Resources
- QA

# Advantages of Using *Spring Data GemFire*

Uses Spring's powerful, non-invasive *programming model* and concepts with Pivotal GemFire to simplify configuration and development

Integration with other powerful *Spring IO Platform* Technologies...



– Spring's *Caching Abstraction*



– Spring's *Transaction Management*



– *Spring Data Commons/REST*



– *Spring Integration* (Inbound/Outbound Channel Adapters)



– *Spring XD* (GemFire Input Source and Output Sink)

And...



**Pivotal GemFire** combined with **Spring Data GemFire** can already be used as a **JSR-107 *caching provider*** when using **Spring Framework 4.1**

# Dependencies Baseline

**Spring Data GemFire (SDG)** was sync'd with the **Spring Data** release train as of *Dijkstra* (~March 2014, v1.4); just released with *Evans* (v1.5)

- *Spring Framework* 4.0.7
- *Spring Data Commons* 1.9
- GemFire 7.0.2
- Built/ran with JDK 7 & JKD 8

# Feature Enhancements

*Region* bean definitions default to create (not lookup)...

```
<gfe:partitioned-region id="Example" .. ignore-if-exists="[true|false]">
  ...
</gfe:partitioned-region>
```

'Manual Start' for *GatewayReceivers*...

```
<gfe:gateway-receiver id="Example" .. manual-start="[true|false]">
  ...
</gfe:gateway-receiver>
```

Ability to execute *Spring Data GemFire* defined GemFire Functions  
using annotations in Gfsh



# Productivity Features

## *Spring Boot Starter Data GemFire (new)...*

```
<dependency>
  <groupId>org.springframework.boot</groupId>
  <artifactId>spring-boot-starter-data-gemfire</artifactId>
</dependency>
```

```
@Configuration
@ImportResource("spring-data-gemfire-cache.xml")
@EnableAutoConfiguration
@EnableGemfireFunctions
@EnableGemfireRepositories
@EnableTransactionManagement
public class SampleDataGemFireApplication {

    public static void main(String[] args) {
        SpringApplication.run(SampleDataGemFireApplication.class, args);
    }
}
```

# Productivity Features

## GemfireTemplate

- simplified GemFire data access...
  - \* `contains[Key|Value]`, `get`, `getAll`, `find`, `query`,
  - \* `put`, `putAll`, `putIfAbsent`, `replace`

## Exception Translation

- maps GemFire Exceptions to `org.springframework.dao` Exceptions
- e.g. `EntryExistsException` → `DuplicateKeyException`

# Productivity Features

Automatic creation of *DiskStore* “Disk Directory Location” (**new**)...

```
<gfe:disk-store id="Example" auto-compact="true" ...>  
  <gfe:disk-dir location="/path/to/disk/files" max-size="1024"/>  
</gfe:disk-store>
```

# Region Lookup 101

Spring Data GemFire supports GemFire `cache.xml`...

```
<gfe:cache cache-xml-location="/path/to/cache.xml" ...>
  ...
</gfe:cache>
```

GemFire 8 Cluster-based Configuration Management...

```
<util:map id="gemfireProperties">
  <prop key="name">SpringGemFireServer</prop>
  <prop key="use-cluster-configuration">[true|false]</prop>
</util:map>

<gfe:cache properties-ref="gemfireProperties" ...>
  ...
</gfe:cache>
```

# Region Lookup

Given GemFire `cache.xml`...

```
<?xml version="1.0"?>
<!DOCTYPE cache PUBLIC "-//GemStone Systems, Inc.
    //GemFire Declarative Caching 7.0//EN"
    "http://www.gemstone.com/dtd/cache7_0.dtd">
<cache>

  <region name="Parent" refid="REPLICATE">
    <region name="Child" refid="REPLICATE">
      <region name="Grandchild" refid="REPLICATE"/>
    </region>
  </region>

  <region name="Example" refid="PARTITION"/>

</cache>
```

Spring `context.xml` Was...

```
<gfe:cache cach-xml-location="/path/to/cache.xml"/>
  <gfe:lookup-region id="Parent">
    <gfe:lookup-region name="Child">
      <gfe:lookup-region name="Grandchild"/>
    </gfe:lookup-region>
  </gfe:lookup-region>

  <!-- or... -->
  <gfe:lookup-region id="/Parent/Child"/>

  <gfe:lookup-region id="Example"/>
```

# Region Lookup

```
@Repository("exampleRepo")
public class RegionBasedRepository extends DaoSupport {

    @Resource(name = "Parent")
    private Region<?, ?> parent;

    @Resource(name = "/Parent/Child")
    private Region<?, ?> child;

    @Resource(name = "/Parent/Child/Grandchild")
    private Region<?, ?> grandchild;

    @Resource(name = "/Example")
    private Region<?, ?> example;

    ...
}
```

# Auto Region Lookup

Given same GemFire `cache.xml`...

```
<?xml version="1.0"?>
<!DOCTYPE cache PUBLIC "-//GemStone Systems, Inc.
    //GemFire Declarative Caching 7.0//EN"
    "http://www.gemstone.com/dtd/cache7_0.dtd">
<cache>

    <region name="Parent" refid="REPLICATE">
        <region name="Child" refid="REPLICATE">
            <region name="Grandchild" refid="REPLICATE"/>
        </region>
    </region>

    <region name="Example" refid="PARTITION"/>

</cache>
```

Spring `context.xml` Now...

```
<gfe:cache cach-xml-location="/path/to/cache.xml"/>

<!--
<gfe:lookup-region id="Parent">
    <gfe:lookup-region name="Child">
        <gfe:lookup-region name="Grandchild"/>
    </gfe:lookup-region>
</gfe:lookup-region>
-->

<gfe:auto-region-lookup/>

<!-- optional -->
<gfe:lookup-region id="Example"/>
```

# Auto Region Lookup

```
@DependsOn("gemfireCache")
@Repository("exampleRepo")
public class RegionBasedRepository extends DaoSupport {

    @Resource(name = "Parent")
    private Region<?, ?> parent;

    @Resource(name = "/Parent/Child")
    private Region<?, ?> child;

    @Resource(name = "/Parent/Child/Grandchild")
    private Region<?, ?> grandchild;

    @Resource(name = "/Example")
    private Region<?, ?> example;

    ...
}
```



# Region Templates

Region Templates are a way to share common configuration across multiple Region bean definitions in the Spring context...

```
<gfe:region-template id="base"/>
```

```
<gfe:replicate-region-template id="replica-tmpl" template="base"/>
```

```
<gfe:partition-region-template id="partition-tmpl" template="base"/>
```

```
<gfe:local-region-template id="local-tmpl" template="base"/>
```

```
<gfe:client-region-template id="client-tmpl" template="base"/>
```

Example...

```
...
<gfe:region-template id="Base" concurrency-checks-enabled="true"
    initial-capacity="51" load-factor="0.85" persistent="false" statistics="true">
    <gfe:cache-listener>
        <bean class="example.CacheListener" p:name="X"/>
    </gfe:cache-listener>
    <gfe:entry-tti timeout="600" action="DESTROY"/>
</gfe:region-template>

<gfe:region-template id="Ext" persistent="true" template="Base">
    <gfe:gateway-sender-ref bean="ExampleGatewaySender"/>
</gfe:region-template>

<gfe:partitioned-region-template id="PartTmpl" copies="2" local-max-memory="8192"
    template="Ext">
    <gfe:cache-loader>
        <bean class="example.CacheLoader"/>
    </gfe:cache-loader>
    <gfe:subscription type="ALL"/>
</gfe:partitioned-region>

<gfe:partitioned-region id="Example" colocated-with="Neighbor" copies="1"
    template="PartTmpl"/>

<gfe:partitioned-region id="Neighbor" persistent="false" template="Ext"/>
...
```

# Region Template Oops

Cannot mix *Region* types; must be compatible...

```
<gfe:partition-region-template id="partition-tmpl" template="base"/>
```

```
<gfe:replicate-region id="Example" template="partition-tmpl"/>
```

# Spring Data *Repositories* for GemFire

Spring Data GemFire's *Repository* is a provider implementation building on Spring Data Common's *Repository Abstraction*...

What is the *Repository Abstraction*...

- CRUD
- Querying
- **Pagination/Slices** & Sorting

Persisting Multiple Regions

# Using Spring Data GemFire *Repositories*

## Domain Class...

```
@Region("Gemstones")
public class Gemstone {

    @Id
    private Long id

    private String name;

    ...
}
```

## Repository...

```
public interface GemstoneRepository extends GemfireRepository<Gemstone, Long> {
    ...
}
```

# Repository Queries

```
public interface GemstoneRepository extends GemfireRepository<Gemstone, Long> {  
  
    Gemstone findByName(String name);  
  
    // @Query("FROM /Gemstones g WHERE g.price > $1")  
    Gemstone findByPriceGreaterThan(BigDecimal price)  
  
    List<Gemstone> findByNameOrderByNameDesc(Sort order);  
}
```

Using *@Query* is useful for *joins* or *object graph traversal*...

```
public interface CustomerRepository extends GemfireRepository<Customer, Long> {  
  
    @Query("<trace> SELECT DISTINCT c FROM /Customers c, c.addresses a  
           WHERE a.city = $1")  
    List<Customers> findCustomersInCity(String city)  
}
```

# @Region Repository

```
@Region("Gemstones")  
public interface GemstoneRepository extends GemfireRepository<Gemstone, Long> {  
}
```

```
@Region("ValuableRocks")  
public interface GemstoneRepository extends GemfireRepository<Gemstone, Long> {  
}
```

# Transaction Management

*Spring Data GemFire* integrates nicely with the *Spring Framework's* **Transaction Management** infrastructure...

- supports both **Cache** and **Global** Transactions

For **Cache** Transactions...

```
<gfe:cache id="cache" .../>
```

```
...
```

```
<gfe:transaction-manager id="txManager" cache-ref="cache"/>
```



# Transaction Management

@Service component...

```
@Service("gemstoneService")
public class GemsService {

    protected static final List<String> APPROVED_GEMS = new ArrayList<String>(
        Arrays.asList( "ALEXANDRITE", "AQUAMARINE", "DIAMOND", "OPAL",
            "PEARL", "RUBY", "SAPPHIRE", "SPINEL", "TOPAZ" ));

    @Transactional(readOnly = false)
    public Gemstone save(Gemstone gemstone) {
        gemstone = getGemFireGemsDao().save(getDatabaseGemsDao().save(gemstone));

        // NOTE deliberate, but stupid and naive business validation
        // after the mutating data access!
        if (!APPROVED_GEMS.contains(gemstone.getName().toUpperCase())) {
            // NOTE if the gemstone is not valid, blow chunks
            // (should cause transaction to rollback for GemFire and Database)!
            throw new IllegalGemstoneException(
                String.format("'%1$s' is not a valid gemstone!",
                    gemstone.getName()));
        }

        return gemstone;
    }
}
```

# GemFire Function Support

Spring Data GemFire provides POJO-based, Function annotation support for creating, registering and executing GemFire Functions

- *@GemfireFunction*
- *@FunctionExecution*

# Creating Function(s)

```
@Component
public class RegionFunctions {

    @GemfireFunction
    public String echo(String message) {
        return String.format("You said '%1$s'!", message);
    }

    @GemfireFunction
    public Integer regionSize(FunctionContext context, String regionNamePath) {
        Region region = getRegion(context, regionNamePath);

        if (region != null) {
            return region.size();
        }

        throw new RegionNotFoundException("The Region on which the size will be determined"
            was not found!");
    }

    protected Region getRegion(FunctionContext context, String regionNamePath) {
        ...
    }
}
```

# Registering Function(s) with GemFire

## XML...

```
<gfe:annotation-driven/>

<bean class="example.RegionFunctions"/>
```

## Component-scanning...

```
<gfe:annotation-driven/>

<context:component-scan base-package="example">
    ...
</context:component-scan>
```

## Java Config...

```
@Configuration
@ImportResource("/spring-data-gemfire-context.xml")
@EnableGemfireFunctions
public class SpringGemFireConfiguration { ... }
```

# Function Execution

## Peer Cache...

```
@OnMember(groups = "exampleGroup")  
public interface OnMemberEchoFunctionExecution {  
  
    String echo(String message);  
}
```

## Client Cache...

```
@OnServer(cache = "gemfireCache")  
public interface OnServerEchoFunctionExecution {  
  
    String echo(String message);  
}
```

# Function Execution

Peer Cache and Client Cache...

```
@OnRegion(region = "Example")
public interface OnRegionFunctionExecution {

    String echo(String message);

    Integer regionSize(String regionNamePath);
}
```

# Function Execution Configuration

XML...

```
<gfe-data:function-executions base-package="example">
  <gfe-data:exclude-filter type="regex" expression=".*OnServer.*"/>
</gfe-data:function-executions>
```

Java Config...

```
@Configuration
@ImportResource("/spring-data-gemfire-context.xml")
@EnableGemfireFunctionExecutions
public class SpringGemFireConfiguration {
    ...
}
```

# Function Execution Invocation

```
@Component
public class ApplicationComponent {

    @Autowired
    private OnRegionFunctionExecution functions;

    public ... someMethod(String regionNamePath) {
        Integer size = functions.regionSize(regionNamePath);
        ...
    }

    ...
}
```



# Order Application Demo

# Spring Data GemFire Case Studies

# Case Study: GemFire Configuration with SDG

Set the stage... 3 Primary Use Cases of *Spring Data GemFire*...

## 1. GemFire Configuration

- replacement for GemFire `cache.xml`

## 2. Peer Cache Application

```
<gfe:cache properties-ref=".." cache-xml-location=".." .../>
```

## 3. Client Cache Application

```
<gfe:client-cache properties-ref=".." cache-xml-location=".." .../>
```

But, how do you “launch” the GemFire Sever?

# Java Main

```
public class SpringGemFireLauncher {  
  
    // e.g. args[0] = "classpath:spring-gemfire-context.xml"  
    private static void validate(String[] args) {  
        ...  
    }  
  
    private static void main(String[] args) {  
        validate(args);  
        new ClassPathXmlApplicationContext(args);  
        ...  
    }  
}
```

# 1.4 Initializer

```
<?xml version="1.0"?>
<!DOCTYPE cache PUBLIC "-//GemStone Systems, Inc.
//GemFire Declarative Caching 7.0//EN"
"http://www.gemstone.com/dtd/cache7_0.dtd">
<cache>
  <initializer>
    <class-name>
org.springframework.data.gemfire.support.SpringContextBootstrappingInitializer
    </class-name>
    <parameter name="contextConfigLocations">
      <string>
classpath:com.example.package,org.example.another.package...
      </string>
    </parameter>
  </initializer>
</cache>
```

# 1.4 Initializer

```
<?xml version="1.0"?>
<!DOCTYPE cache PUBLIC "-//GemStone Systems, Inc.
//GemFire Declarative Caching 7.0//EN"
"http://www.gemstone.com/dtd/cache7_0.dtd">
<cache>
  <region name="Example" refid="REPLICATE">
    <region-attributes initial-capacity="51" load-factor="0.85">
      <cache-loader>
        <class-name>example.CacheLoader</class-name>
      </cache-loader>
    </region-attributes>
  </region>
  <initializer>
    <class-name>
org.springframework.data.gemfire.support.SpringContextBootstrappingInitializer
    </class-name>
    <parameter name="contextConfigLocations">
      <string>
classpath:path/to/spring-gemfire-context.xml,
classpath:path/to/app-context.xml
      </string>
    </parameter>
  </initializer>
</cache>
```

# 1.4 LazyWiringDeclarableSupport

Used to be, use *Spring Data GemFire's* WiringDeclarableSupport

```
<!-- Spring context.xml -->
<gfe:cache cache-xml-location="/path/to/cache.xml" ... />

<!-- GemFire cache.xml -->
<cache>
  <region name="Example" refid="REPLICATE">
    <region-attributes initial-capacity="51" load-factor="0.85">
      <cache-loader>
        <class-name>example.CacheLoader</class-name>
```

Now, must use...

```
package example;
public class CacheLoader
    extends LazyWiringDeclarableSupport
    implements CacheLoader<String, User> {

    @Autowired
    private DataSource dataSource;

    ...
}
```



# GemFire 8 Gfsh Support

```
gfsh>start server -spring-xml-location="classpath:path/to/context.xml"
```

*Resource* Location Prefixes...

```
-spring-xml-location="file:/path/to/context.xml"
```

# DEMO

# Case Study: Global Transaction Management

Using GemFire with an external data store (e.g. MySQL RDBMS) in a Global Transactional context...

- Use Spring's `JtaTransactionManager`; Not...

```
<!-- WRONG -->
```

```
<gfe:transaction-manager id="txManager" cache-ref="cache"/>
```

- Uses **GemFire's JTA Transaction Manager** provider implementation along with GemFire's **JNDI context** to *register DataSource* as well as *locate GemFire's TX Manager and JTA UserTransaction*
- Why not **Atomikos, Bitronix, or JOTM**?

Code...

# DEMO

# Case Study: Spring Caching Abstraction

Use Spring's *Caching Abstraction*...

Use GemFire as ***caching provider*** to replace ***Oracle Coherence***, but absolutely no presence of GemFire in the application...

Support for varying (some old, some new) application client ***versions***...

- New Domain Object class field “additions”
- Old Domain Objects should trigger a “cache miss”
- New Domain Object state must be preserved when used by an old application (PDX)

The **10% case** *Juergen* was talking about in *Caching and Messaging Improvements*

# DEMO

# Looking Ahead...

P0: Finish **GemFire 8** Support in **Spring Data GemFire 2.0!**

P0: Edit ***Spring Data GemFire Reference Guide***

P0: Rewrite ***Spring GemFire Examples***

P0: Provide *Spring GemFire Case Studies*

P0: Provide *Spring Data GemFire Reference Implementation*

P1: Integration Test Framework for Distributed Test Cases, WAN

P1: Integration with *Spring Security* as a facade over GemFire's Security Model

P2: Support for *Java-based Configuration* (Groovy)

P3: *Pagination* and *Slices*

# Resources

GemFire Overview

[http://gemfire.docs.pivotal.io/latest/userguide/index.html#getting\\_started/product\\_intro.html](http://gemfire.docs.pivotal.io/latest/userguide/index.html#getting_started/product_intro.html)

GemFire User Guide

<http://gemfire.docs.pivotal.io/latest/userguide/index.html>

Spring Data GemFire Reference Guide

<http://docs.spring.io/spring-data-gemfire/docs/current/reference/html/>

Guide: <https://spring.io/guides/gs/accessing-data-gemfire/>

Guide: <https://spring.io/guides/gs/caching-gemfire/>

GitHub (source): <https://github.com/spring-projects/spring-data-gemfire>

GitHub (examples): <https://github.com/spring-projects/spring-gemfire-examples>

JIRA: <https://jira.spring.io/browse/SGF>

StackOverflow: <https://stackoverflow.com/questions/tagged/spring-data-gemfire>



Thank You!